


Graph Algorithms

Note Title

9/25/2014

Representation



Graph has n nodes m edges

Adjacency Matrix

	1	2	3	4
1	0	1	0	0
2	0	0	1	1
3	0	0	0	1
4	0	0	0	0

directed

	1	2	3	4
1	0	1	0	1
2	1	0	1	1
3	0	1	0	1
4	1	1	1	0

undirected

Storage alternatives

- ~ store matrix on disk (each row is line)
- $n^2 \times n$

most graphs are sparse

- store coordinates of non zero entries.
 $2m$

- Adjacency list

node	node list of neighbors
1	2 y
2	3 y $m+n$
3	y

Shortest path Algorithm

Dijkstra's alg. Find shortest path between source and all other nodes in

weighted graph / $w(v, v) \geq 0$

Dijkstra(G, w, s) = {

$d[s] = 0$

For every other node v :

$d[v] = \text{infinity}$ (or $n+1$)

Priority $Q = \{s\}$ (priority based on shortest d value)

while(Q is not empty) {

$v \leftarrow Q.\text{getmin}()$

for every out neighbor u of v :

if $d[u] > d[v] + w(v, u)$ then

$d[u] = d[v] + w(v, u)$

backpointer[u] = v (optional)


$Q.\text{add}(u)$

Mapper

```
map(key: nodeid, value="(nid, weight) (nid, weight) .... d[nodeid]") {  
    dist = d[nodeid]  
    emit(key, value) //structure of graph passed to reducer  
    for each node m in value:  
        emit(m, dist + weight(nodeid, m))  
}
```

Reducer

```
reduce(key: nodeid, valuelist = {adjlist, count1, count2, count3, ...})  
    dmin = infinity  
    for each value:  
        if value is adjlist  
            M <- value  
        else //value is a number  
            if val < dmin:  
                dmin <- value  
    if M.distance < dmin:  
        dmin = M.distance  
    if dmin < M.distance: increment counter  
    M.distance = dmin  
    emit(nodeid, M)
```



Random surfer

with prob α = goto random web page

$(1-\alpha)$ = pick one of outgoing links from current page and follow it

importance = fraction of time surfer spends

$$P_r(V) = \frac{\alpha}{n} + (1-\alpha) \sum_{u \text{ in links to } V} \frac{P_r(u)}{\text{outdegree}(u)} \quad \left| \begin{array}{l} \text{normalize} \\ \text{so} \\ \text{sums to } 1 \end{array} \right.$$

Mapper

```
Map(key: nodeid, value = "n1 n2 n3 pr") {  
  p <- pr / number of nodes nodeid points to  
  emit(nodeid, value)  
  for every node m that nodeid points to:  
    emit(m, p)
```

To add teleportation and reweighting so that pageranks sum up to 1, do the following

1. mapper multiplies p by $1-\alpha$
2. in reducer, inp is initialized with α/n where n is the number of total nodes in the graph.
3. Mapper maintains sum of the pageranks it emits. these sums get sent to all reducers.

Reducer

```
reduce(key: node id, value-list = {adj-list, p1 p2 p3 p4})  
  inp = 0  
  for each value in value-list:  
    if value is adj-list  
      M <- value  
    else  
      inp += value  
  M.pr = inp  
  emit(nodeid, M)
```

4. Each reducer aggregates the special values from step 3 and adds alpha to them (this must be done first using the order inversion technique). This sum is the normalizing factor

5. Before emitting, the reducer divides inp by the normalizing factor.