

Stochastic Gradient Descent

CSE *97B

Big Data



Empirical Risk

- Data = $\{(\vec{X}_1, t_1), (\vec{X}_2, t_2), \dots, (\vec{X}_n, t_n)\}$
 - \vec{X}_i is the feature vector (set of measurements) for the i^{th} record:
 $\vec{X}_i = (X_i[0], X_i[1], \dots, X_i[m])$.
 - t_i is the target for the i^{th} record
- Models have parameters (e.g. θ, \vec{w}, b)
 - Linear regression: $\hat{t} = \vec{w} \cdot \vec{X} + b$
 - Logistic regression: $\hat{P}(t) = \frac{1}{1 + \exp(-t(\vec{w} \cdot \vec{X} + b))}$
 - $t = \pm 1$ (e.g. $t = 1$ if rain, $t = -1$ if no rain).
 - Logistic regression models the probability of an event.
- Parameters set by optimization problem (minimize empirical risk).
 - Linear regression: $\vec{w}, b = \operatorname{argmin} \sum_{i=1}^n (\vec{w} \cdot \vec{X}_i + b - t_i)^2$
 - Logistic regression: $\vec{w}, b = \operatorname{argmin} \sum_{i=1}^n \log(1 + \exp(-t_i(\vec{w} \cdot \vec{X}_i + b)))$
 - Common pattern: $\vec{w}, b = \operatorname{argmin} \sum_{i=1}^n f(\vec{X}_i, t_i, \vec{w}, b)$



Minimizing Empirical Risk

- Linear regression: $\vec{w}, b = \operatorname{argmin} \sum_{i=1}^n (\vec{w} \cdot \vec{X}_i + b - t_i)^2$
 - Empirical risk $R(\vec{w}, b) = \sum_{i=1}^n (\vec{w} \cdot \vec{X}_i + b - t_i)^2$
 - Can be solved in closed form (i.e. formula).
 - Compute $\frac{\partial R}{\partial w[0]}, \frac{\partial R}{\partial w[1]}, \dots, \frac{\partial R}{\partial w[m]}, \frac{\partial R}{\partial b}$
 - Set partial derivatives to 0 and solve.
- Logistic regression: $\vec{w}, b = \operatorname{argmin} \sum_{i=1}^n \log \left(1 + \exp(-t_i(\vec{w} \cdot \vec{X}_i + b)) \right)$
 - No closed form solution, need algorithm.
- Common pattern: $\vec{w}, b = \operatorname{argmin} \sum_{i=1}^n f(\vec{X}_i, t_i, \vec{w}, b)$
 - Generally no closed form solution, need algorithm.
- mini-batch Stochastic gradient descent: simple and fast.



Mini-Batch Stochastic Gradient Descent

Goal: find \vec{w}, b to minimize $\sum_{i=1}^n f(\vec{X}_i, t_i, \vec{w}, b)$

parameters: Batch size k ; learning rate η // Typically $k = 10$

Initialize b and \vec{w}

while *stopping criterion not met* **do**

Select next k records $(\vec{X}_j, t_j), (\vec{X}_{j+1}, t_{j+1}), \dots, (\vec{X}_{j+k-1}, t_{j+k-1})$

// Error contribution is $\sum_{i=j}^{j+k-1} f(\vec{X}_i, t_i, \vec{w}, b)$

$b \leftarrow b - \eta \sum_{i=j}^{j+k-1} \frac{\partial f(\vec{X}_i, t_i, \vec{w}, b)}{\partial b}$

for $\ell = 0..m$ **do**

$w[\ell] \leftarrow w[\ell] - \eta \sum_{i=j}^{j+k-1} \frac{\partial f(\vec{X}_i, t_i, \vec{w}, b)}{\partial w[\ell]}$

end

if *at end of dataset* **then**

Permute dataset;

end

end



In MapReduce

- Algorithm is inherently sequential.
- Cannot be parallelized as is.
- Can be approximated in parallel mode.
 - Zinkevich et al. “Parallelized Stochastic Gradient Descent” NIPS 2010.
- Main idea:
 - Each mapper independently performs stochastic gradient descent.
 - Same initial \vec{w} and b
 - A reducer averages the \vec{w} from all mappers (similarly for b)
 - This is the common starting point for the next iteration.
- We use sort and shuffle phase to permute the data.



Controller

```
// For simplicity, batch size  $k = 1$   
Set learning rate  $\eta$  for each mapper  
Initialize  $\vec{w}$  and  $b$ , distribute to all mappers  
while convergence criterion not met do  
    Set up mapreduce job  
    Obtain  $\vec{w}$  and  $b$  from Reducer 0 // via file or Hive  
    Distribute  $\vec{w}$  and  $b$  to each mapper  
    Modify  $\eta$  if necessary, set for each mapper  
end
```



Mapper

def *setup()*:

- | Get parameters η
- | Read initial \vec{w} and b

def *map(key, value)*:

- | $(\vec{X}, t) = \text{parse}(\text{value})$
- | $\text{emit}(\text{RandomPositiveNumber}, \text{value})$ // for permuting data
- | $b \leftarrow b - \eta \frac{\partial f(\vec{X}, t, \vec{w}, b)}{\partial b}$
- | **for** $\ell = 0..m$ **do**
 - | $w[\ell] \leftarrow w[\ell] - \eta \frac{\partial f(\vec{X}, t, \vec{w}, b)}{\partial w[\ell]}$
- | **end**

def *cleanup()*:

- | $\text{emit}(\vec{w} \text{ and } b \text{ to Reducer } 0)$



Reducer

```
def reduce(key, value-list):  
    if value-list is the set of vectors:  $\vec{w}_1, \vec{w}_2, \dots, \vec{w}_\ell$  then  
        |  $\vec{w} = \frac{1}{\ell} \sum_{i=1}^{\ell} \vec{w}_i$   
    else if value-list is the set of  $b$  values:  $b_1, b_2, \dots, b_\ell$  then  
        |  $b = \frac{1}{\ell} \sum_{i=1}^{\ell} b_i$   
    else  
        | // Ignore key (which is random number)  
        | // value is data record, output it  
        | For every value in value-list, emit(value)  
    end  
  
def cleanup():  
    | Output  $\vec{w}$  and  $b$  (e.g. to HDFS or Hive)
```

